

You can't handle the lie: Catching lying actors in the BGP protocol

Clay Thomas
claytont@cs.princeton.edu

Gavriel Hirsch
gbhirsch@cs.princeton.edu

November 27, 2018

In many BGP networks, nodes can get their preferred path by advertising a certain path, but forwarding traffic along another. However, in many of the example networks from the literature, other nodes can collectively detect these lies by comparing the route that is advertised (known by the node the manipulator lies to) to the route actually taken (known by the node the manipulator actually forwards traffic to). We want to study the conditions under which nodes can lie without being detected, as well as the economic incentives of other nodes to collaboratively check each other (such as providers helping customers detect liars).

Our thinking is primarily based on the routing games model of [LSZ08], where node's objectives are simply to get better paths to a destination. Additional consideration is given to the model of [GHJ⁺08], where nodes seek to attract traffic as well as get better paths.

1 In the [LSZ08] model, you can always catch a liar

Intuitively, if a node lies to get a better path, the node it lies to and the node it routes through must be connected. The following conjecture, which we hope to prove formally, makes this precise.

Conjecture 1. *Suppose No Dispute Wheel holds, but route verification does not, and assume that the network is connected. Suppose that (assuming other nodes play truthfully) a node m can achieve a better path to d by announcing a route that does not exist to a node v . Let m 's next hop in the manipulated routing tree be denoted r . Then there exists a path in the network, not containing m , between v and r .*


This means that, if all nodes other than m are fully collaborative and honest, the nodes will be able to detect m 's lie by communicating along the links that already exist in the network.

(a) In the [GHJ⁺08] model, you can't

Consider Bowtie (figure 1), an example taken directly from [GHJ⁺08]. In this case, node m needs to lie to nodes c and n , but actually forward traffic to node d . The only paths between

d and c (or n) pass through node m . Thus, the honest nodes cannot catch m in its lie solely using the communication channels provided by the network itself. With this example, we see that the [GHJ⁺08] model provides more complications than [LSZ08] for our questions.


Figure 1: Bowtie



2 But will you?

Figure 2 shows an example network (taken directly from [LSZ08] but rendered in the style of [GHJ⁺08]) where m has an incentive to advertise a path that does not exist. We regard this network as the canonical example of Conjecture 1. Indeed, node 1 can inform node 2 that it is receiving traffic from m , so 2 will know that m is lying about its path.

Figure 2: Nonexistent



This example is a Gao-Rexford network, where customer-provider relationships are denoted by arrows pointing from customers to providers. Note that node 2 must be a provider of node 1 for this example (otherwise, 2 would not export its route to 1). In general, customers do not provide services for their providers, so it may not make sense for node 1 to help node 2 catch m 's lie. On the other hand, if 1 does inform 2 of the lie, and 2 resumes using its old route, 1 actually gets a path that it prefers.

Given the above considerations, we regard the competing incentives of

- getting good paths
- economic obligations of customer-provider relationships
- costs of maintaining and communicating lie-detecting information

to be an interesting avenue of research.

3 Project Components

We plan to first analyze which situations can result in a node having an incentive to lie and other nodes having incentives to not expose that lie to each other. In addition we will design a new protocol and consider corresponding business relationships in which nodes can both send messages about their paths, and separately share other information that can be used to expose others' lies. We will also analyze situations under this new protocol and see how the analysis differs. Throughout the process, we hope to state and prove formal statements such as Conjecture 1, in order to gain insights into the strategic nature of BGP.

References

- [GHJ⁺08] Sharon Goldberg, Shai Halevi, Aaron D. Jaggard, Vijay Ramachandran, and Rebecca N. Wright. Rationality and traffic attraction: Incentives for honest path announcements in bgp. In *Proceedings of the ACM SIGCOMM 2008 Conference on Data Communication*, SIGCOMM '08, pages 267–278, New York, NY, USA, 2008. ACM.
- [LSZ08] Hagay Levin, Michael Schapira, and Aviv Zohar. Interdomain routing and games. In *Proceedings of the Fortieth Annual ACM Symposium on Theory of Computing*, STOC '08, pages 57–66, New York, NY, USA, 2008. ACM.